

Apache Flink hands-on



instructions

Vasiliki Kalavri
kalavriv@inf.ethz.ch
[@vkalavri](#)

Resources

- **Training:** <http://training.data-artisans.com/>
- **Setup instructions** (own laptop): <http://training.data-artisans.com/devEnvSetup.html>
- **Data:** <http://training.data-artisans.com/exercises/taxiData.html>
- **DataStream API documentation:** https://ci.apache.org/projects/flink/flink-docs-release-1.6/dev/datastream_api.html

Taxi Rides

rideId	: Long	// a unique id for each ride
taxiId	: Long	// a unique id for each taxi
driverId	: Long	// a unique id for each driver
isStart	: Boolean	// TRUE for start events, FALSE for end
startTime	: DateTime	// the start time of a ride
endTime	: DateTime	// the end time of a ride, // "1970-01-01 00:00:00" for start
startLon	: Float	// the longitude of the start location
startLat	: Float	// the latitude of the start location
endLon	: Float	// the longitude of the end location
endLat	: Float	// the latitude of the end location
passengerCnt	: Short	// number of passengers on the ride

Taxi Fares

rideId	: Long	// a unique id for each ride
taxiId	: Long	// a unique id for each taxi
driverId	: Long	// a unique id for each driver
startTime	: DateTime	// the start time of a ride
paymentType	: String	// CSH or CRD
tip	: Float	// tip for this ride
tolls	: Float	// tolls for this ride
totalFare	: Float	// total fare collected

Generate Ride Events

```
// get an ExecutionEnvironment
StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

// configure event-time processing
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

// get the taxi ride data stream
DataStream<TaxiRide> rides = env.addSource(
    new TaxiRideSource("/path/to/nycTaxiRides.gz", maxDelay,
    servingSpeed);
```

max event delay

event speedup factor

Generate Fare Events

```
// get an ExecutionEnvironment
StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

// configure event-time processing
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

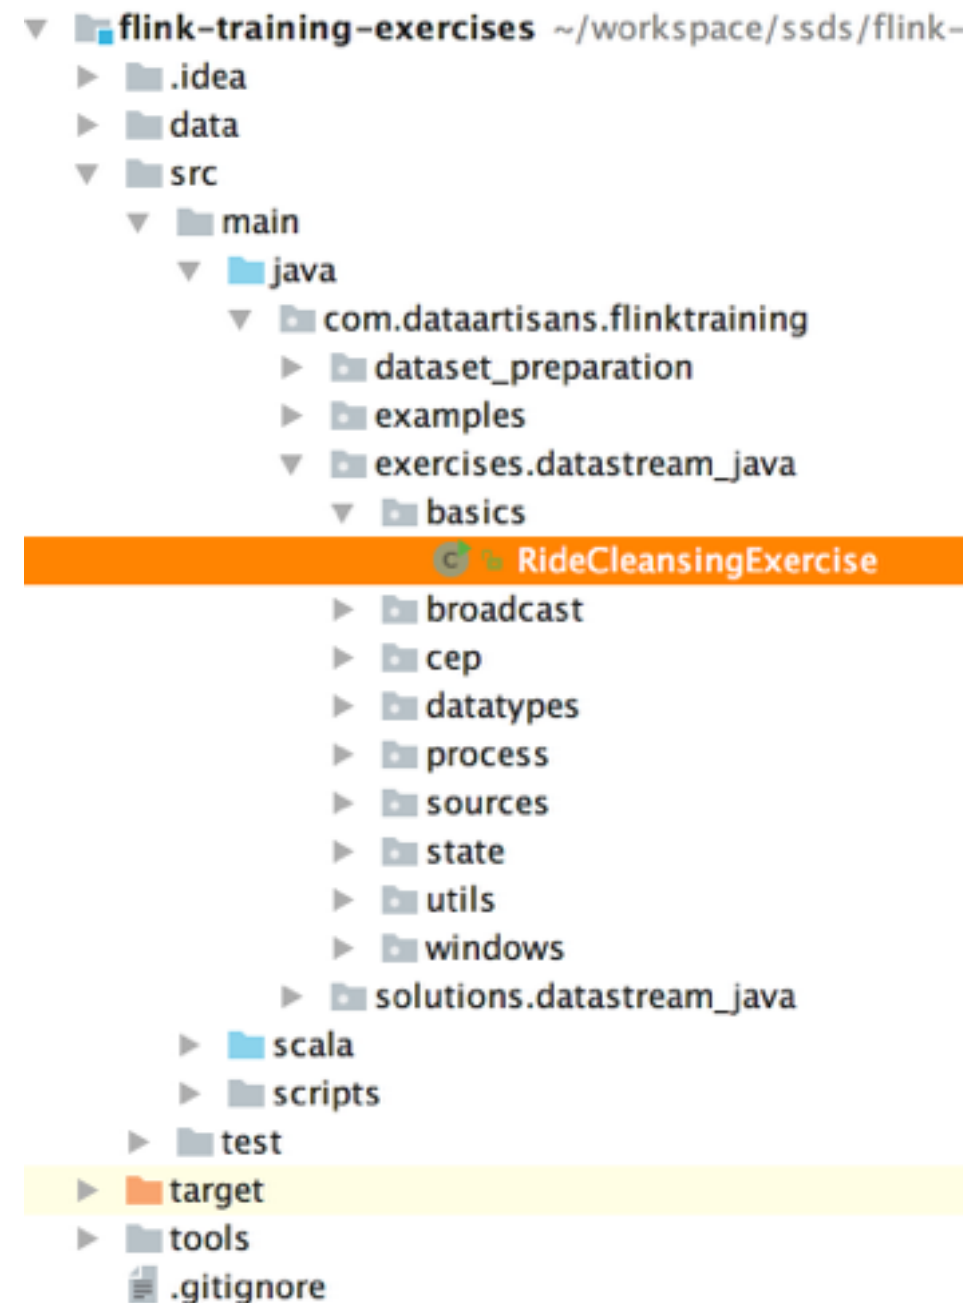
// get the taxi fare data stream
DataStream<TaxiFare> rides = env.addSource(
    new TaxiFareSource("/path/to/nycTaxiFares.gz", maxDelay,
        servingSpeed));
```

Test Setup

1. Open `com.dataartisans.flinktraining.exercises.datastream_java.utils.ExerciseBase` in your IDE
2. Update `pathToRideData` and `pathToFareData`
3. Open `com.dataartisans.flinktraining.examples.datastream_java.basics.RideCount` in your IDE
4. Run the `main()` method
5. Watch the result stream!

Exercise #1: RideCleansing

Filter out rides that do not start/end in NYC



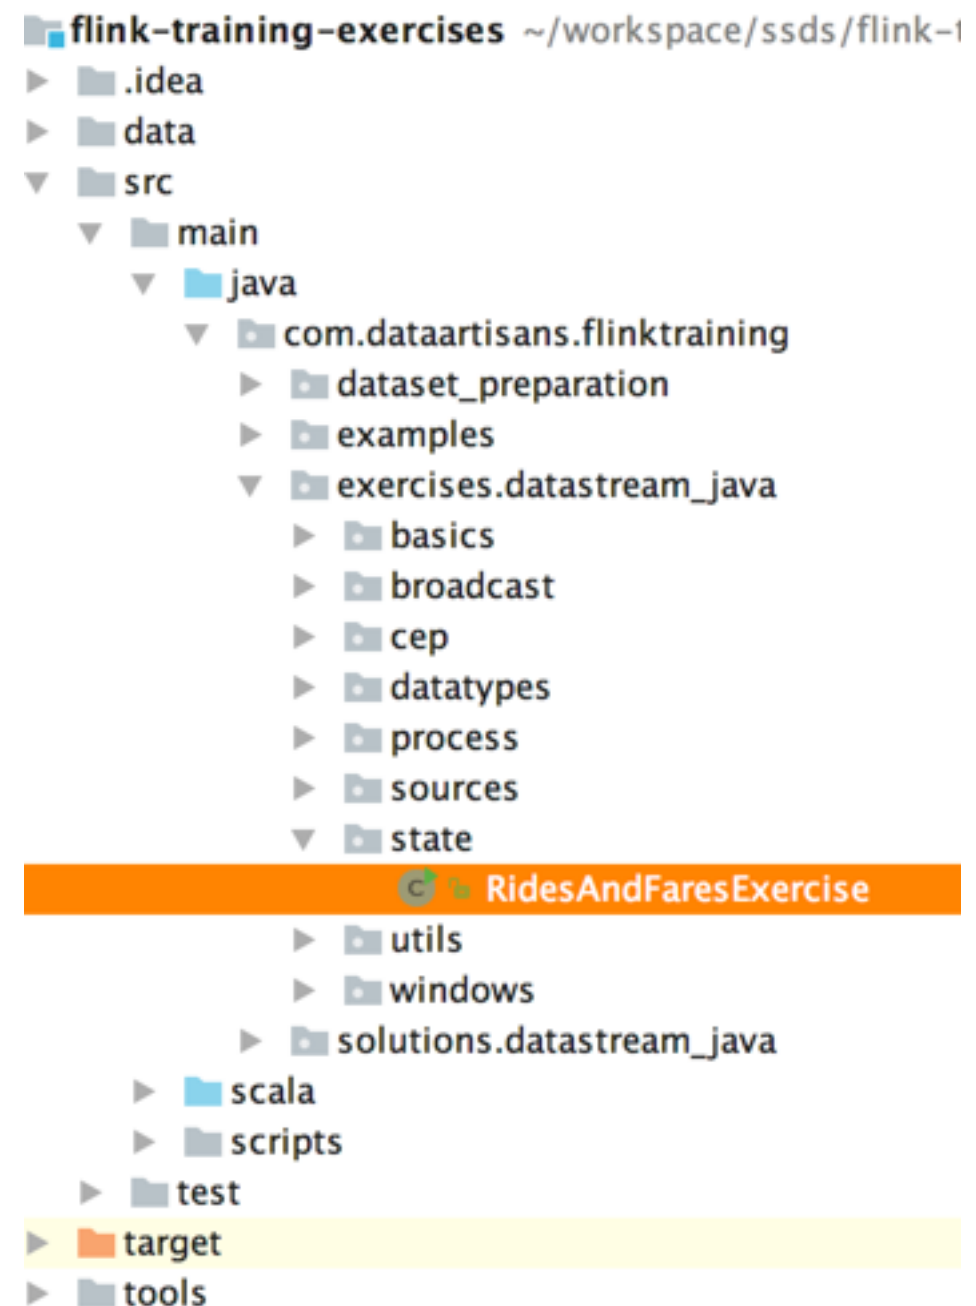
Exercise #1: GeoUtils

```
flink-training-exercises ~/workspace/ssds/flink-trai
├── .idea
├── data
├── src
│   ├── main
│   │   ├── java
│   │   │   ├── com.dataartisans.flinktraining
│   │   │   │   ├── dataset_preparation
│   │   │   │   ├── examples
│   │   │   │   └── exercises.datastream_java
│   │   │   │       ├── basics
│   │   │   │       │   ├── RideCleansingExercise
│   │   │   │       │   ├── broadcast
│   │   │   │       │   ├── cep
│   │   │   │       │   ├── datatypes
│   │   │   │       │   ├── process
│   │   │   │       │   ├── sources
│   │   │   │       │   ├── state
│   │   │   │       │   └── utils
│   │   │   │       │       ├── influxdb
│   │   │   │       │       ├── ConnectedCarAssigner
│   │   │   │       │       ├── ExerciseBase
│   │   │   │       │       ├── GeoUtils
│   │   │   │       │       ├── MissingSolutionException
│   │   │   │       │       ├── TaxiRideSchema
│   │   │   │       │       └── TravelTimePredictionModel
│   │   │   │       ├── windows
│   │   │   │       └── solutions.datastream_java
│   │   ├── scala
│   │   ├── scripts
│   │   └── test
```

```
/**
 * Checks if a location specified by longitude and latitude values is
 * within the geo boundaries of New York City.
 *
 * @param lon longitude of the location to check
 * @param lat latitude of the location to check
 *
 * @return true if the location is within NYC boundaries, otherwise false.
 */
public static boolean isInNYC(float lon, float lat){}
```

Exercise #2: State

Join each TaxiRide
with its TaxiFare

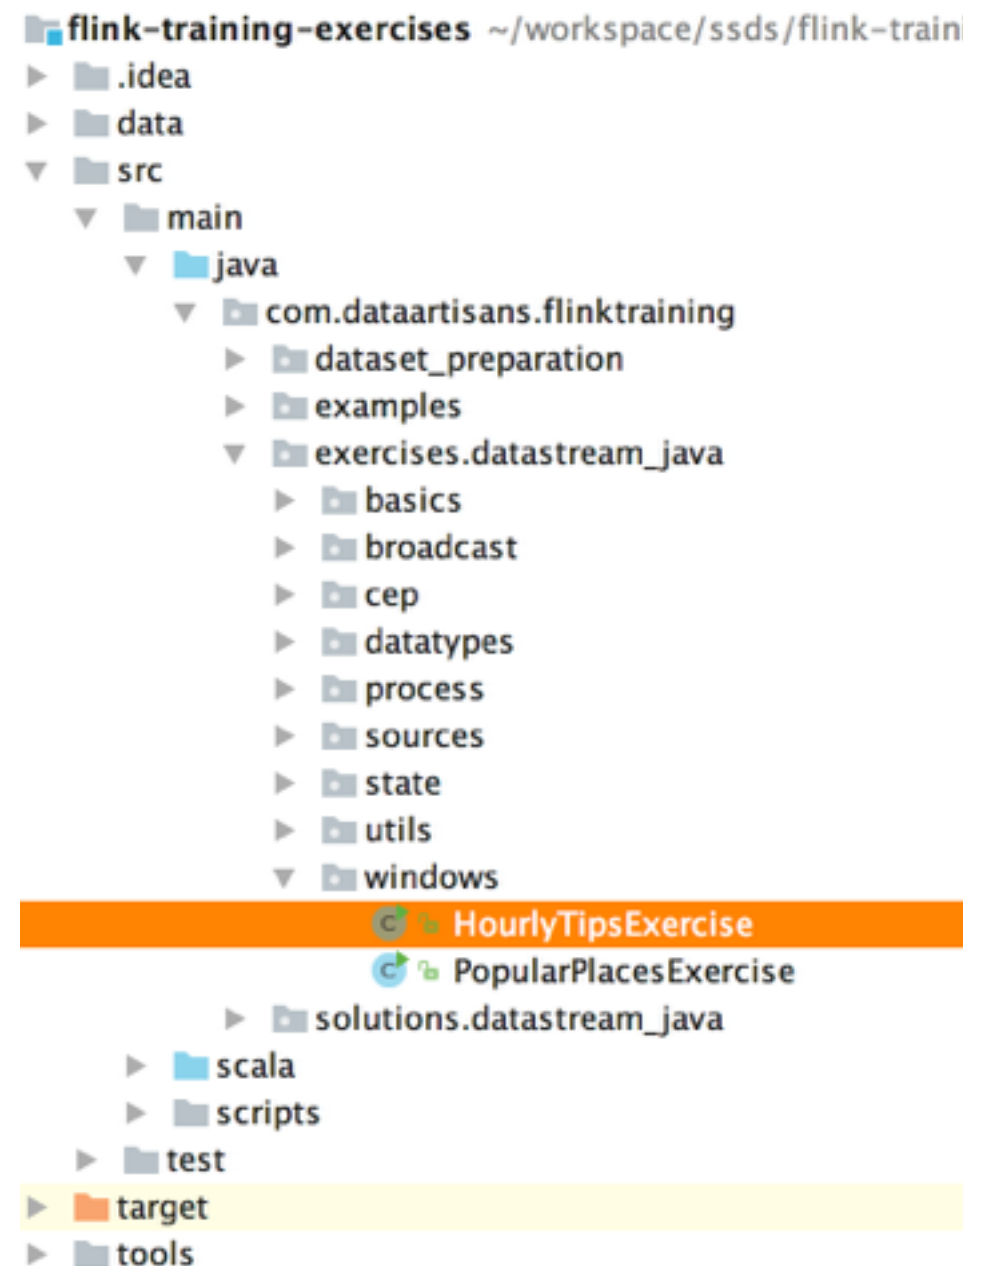


Exercise #2: Hints

- The order of arrival of the ride and fare events is not guaranteed.
- If you receive a fare event, store it until you receive the matching ride and vice versa. Once you have a match, you can emit the result.
- Clear the state once it is no longer needed!
- Check: <https://ci.apache.org/projects/flink/flink-docs-release-1.6/dev/stream/state/state.html#using-managed-keyed-state>

Exercise #3: Windows

Compute which driver
is earning the most
tips every hour



Exercise #3: Hints

1. Compute in 2 steps:
 - 1.1. Total tips per hour per driver
 - 1.2. Driver with max tip
2. Use an incremental `AggregateFunction` with a `ProcessWindowFunction` for 1.1
 - 2.1. <https://ci.apache.org/projects/flink/flink-docs-release-1.6/dev/stream/operators/windows.html#processwindowfunction-with-incremental-aggregation>
3. Use a global window for 1.2